


Experiment Information


Figure:	Regret-notion comparison
Environment:	Historical-frequency
Actions:	10
Feedback:	Bandit feedback
Algorithm:	EXP3
Horizon:	1,000,000
Replicates:	12
Seed:	42
Regrets:	External, Internal, Swap
View:	R / T
Final endpoints at T:	1,000,000
Endpoint statistics:	replicate mean, sample SD (ddof=1), SE=SD/sqrt(n)
External:	mean = 0.004803, SD = 0.00170114, SE = 0.000491078
Internal:	mean = 0.00562192, SD = 0.00148798, SE = 0.000429542
Swap:	mean = 0.0346296, SD = 0.00799226, SE = 0.00230717



- External regret
- Internal regret
- Swap regret


Experiment Information

Figure:	Regret-notion comparison
Environment:	Historical-frequency
Actions:	10
Feedback:	Bandit feedback
Algorithm:	EXP3
Horizon:	1,000,000
Replicates:	12
Seed:	42
Regrets:	External, Internal, Swap
View:	$R / \sqrt{T}$
Final endpoints at T:	1,000,000
Endpoint statistics:	replicate mean, sample SD (ddof=1), SE=SD/sqrt(n)
External:	mean = 4.803, SD = 1.70114, SE = 0.491078
Internal:	mean = 5.62192, SD = 1.48798, SE = 0.429542
Swap:	mean = 34.6296, SD = 7.99226, SE = 2.30717



Experiment Information


Figure:	Regret-notion comparison
Environment:	Historical-frequency
Actions:	10
Feedback:	Bandit feedback
Algorithm:	BM-EXP3
Horizon:	1,000,000
Replicates:	12
Seed:	42
Regrets:	External, Internal, Swap
View:	R / T
Final endpoints at T:	1,000,000
Endpoint statistics:	replicate mean, sample SD (ddof=1), SE=SD/sqrt(n)
External:	mean = 0.00603992, SD = 0.00179243, SE = 0.000517429
Internal:	mean = 0.00199142, SD = 0.000436228, SE = 0.000125928
Swap:	mean = 0.0113095, SD = 0.00147638, SE = 0.000426195



- External regret
- Internal regret
- Swap regret


Experiment Information

Figure:	Regret-notion comparison
Environment:	Historical-frequency
Actions:	10
Feedback:	Bandit feedback
Algorithm:	BM-EXP3
Horizon:	1,000,000
Replicates:	12
Seed:	42
Regrets:	External, Internal, Swap
View:	$R / \sqrt{T}$
Final endpoints at T:	1,000,000
Endpoint statistics:	replicate mean, sample SD (ddof=1), SE=SD/sqrt(n)
External:	mean = 6.03992, SD = 1.79243, SE = 0.517429
Internal:	mean = 1.99142, SD = 0.436228, SE = 0.125928
Swap:	mean = 11.3095, SD = 1.47638, SE = 0.426195



Experiment Information


Figure:	Regret-notion comparison
Environment:	Historical-frequency
Actions:	10
Feedback:	Bandit feedback
Algorithm:	EXP3-IX
Horizon:	1,000,000
Replicates:	12
Seed:	42
Regrets:	External, Internal, Swap
View:	R / T
Final endpoints at T:	1,000,000
Endpoint statistics:	replicate mean, sample SD (ddof=1), SE=SD/sqrt(n)
External:	mean = 0.00419725, SD = 0.0025064, SE = 0.000723535
Internal:	mean = 0.00560775, SD = 0.00108382, SE = 0.000312871
Swap:	mean = 0.036513, SD = 0.00761423, SE = 0.00219804



- External regret
- Internal regret
- Swap regret


Experiment Information

Figure:	Regret-notion comparison
Environment:	Historical-frequency
Actions:	10
Feedback:	Bandit feedback
Algorithm:	EXP3-IX
Horizon:	1,000,000
Replicates:	12
Seed:	42
Regrets:	External, Internal, Swap
View:	$R / \sqrt{T}$
Final endpoints at T:	1,000,000
Endpoint statistics:	replicate mean, sample SD (ddof=1), SE=SD/sqrt(n)
External:	mean = 4.19725, SD = 2.5064, SE = 0.723535
Internal:	mean = 5.60775, SD = 1.08382, SE = 0.312871
Swap:	mean = 36.513, SD = 7.61423, SE = 2.19804




Experiment Information

Figure:	Regret-notion comparison
Environment:	Historical-frequency
Actions:	10
Feedback:	Bandit feedback
Algorithm:	Ito-Tsallis
Horizon:	1,000,000
Replicates:	12
Seed:	42
Regrets:	External, Internal, Swap
View:	R / T
Final endpoints at T:	1,000,000
Endpoint statistics:	replicate mean, sample SD (ddof=1), SE=SD/sqrt(n)
External:	mean = 0.0042845, SD = 0.00121424, SE = 0.000350521
Internal:	mean = 0.00225817, SD = 0.000532163, SE = 0.000153622
Swap:	mean = 0.0124485, SD = 0.00251438, SE = 0.000725839




Experiment Information

Figure:	Regret-notion comparison
Environment:	Historical-frequency
Actions:	10
Feedback:	Bandit feedback
Algorithm:	Ito-Tsallis
Horizon:	1,000,000
Replicates:	12
Seed:	42
Regrets:	External, Internal, Swap
View:	$R / \sqrt{T}$
Final endpoints at T:	1,000,000
Endpoint statistics:	replicate mean, sample SD (ddof=1), SE=SD/sqrt(n)
External:	mean = 4.2845, SD = 1.21424, SE = 0.350521
Internal:	mean = 2.25817, SD = 0.532163, SE = 0.153622
Swap:	mean = 12.4485, SD = 2.51438, SE = 0.725839



Experiment Information

Figure: Regret-notion comparison
Environment: Historical-frequency
Actions: 10
Feedback: Bandit feedback
Algorithm: LCE-IX
Horizon: 1,000,000
Replicates: 12
Seed: 42
Regrets: External, Internal, Swap
View: R / T
Final endpoints at T: 1,000,000
Endpoint statistics: replicate mean, sample SD (ddof=1), SE=SD/sqrt(n)
External: mean = 0.000834583, SD = 0.00119436, SE = 0.000344781
Internal: mean = 0.000978167, SD = 0.000230327, SE = 6.64896e-05
Swap: mean = 0.00415367, SD = 0.00111991, SE = 0.000323289



- External regret
- Internal regret
- Swap regret

Experiment Information

Figure:	Regret-notion comparison
Environment:	Historical-frequency
Actions:	10
Feedback:	Bandit feedback
Algorithm:	LCE-IX
Horizon:	1,000,000
Replicates:	12
Seed:	42
Regrets:	External, Internal, Swap
View:	$R / \sqrt{T}$
Final endpoints at T:	1,000,000
Endpoint statistics:	replicate mean, sample SD (ddof=1), SE=SD/sqrt(n)
External:	mean = 0.834583, SD = 1.19436, SE = 0.344781
Internal:	mean = 0.978167, SD = 0.230327, SE = 0.0664896
Swap:	mean = 4.15367, SD = 1.11991, SE = 0.323289

